Text Segmentation Using Generative Language Models

Pierre Lardet Maya Shah

Abstract

Topic segmentation using Generative Language Models (GLMs) remains relatively unexplored. Previous methods use lexical or semantic similarity between parts of a document to decide on boundaries, but they lack the long range dependency and vast knowledge contained in GLMs. Here, we propose a new prompting strategy and compare to semantic similarity-based methods. Results show that GLMs can be more effective segmenters than existing methods, but issues remain to be solved before GLMs can be relied upon for topic segmentation.

1. Introduction

1.1. The Topic Segmentation Problem

Topic segmentation is the problem of dividing a string of text into constituent parts or 'segments'. Each segment should be semantically self-contained such that it is about one thing. The precise definition of a segment should be dependent upon the specific use cases. For this work, segment boundaries shall always lie on sentence boundaries.

We can think of this as a binary classification task. Given a list of input sentences S, of length n, the model must decide whether there exists a segment boundary between each pair of adjacent sentences. There are $n - 1$ possible boundaries, and therefore our solution space is 2^{n-1}. Formally, the model must find a mapping f from the list of sentences to a binary vector of length $n - 1$: $f(S) = \mathbf{y}$ where $\mathbf{y} = \{y_1, y_2 \ldots, y_{n+1}\}$ and each of the $y_i \in \{0, 1\}$.

Relative to generative tasks such as summarisation, the space of possible solutions is much smaller, but the problem remains subjective as where a boundary should lie can be ambiguous. Frequently, humans cannot agree on a correct solution (Hearst, 1997).

This task can be important as a processing step before some other NLP task, or can be important in its own right. One might use segments to generate a contents page for a long document, or individual summaries of segments within a document, or generate titles for each segment. Other tasks such as information retrieval, long-document summarisation and classification, can all benefit from first being broken down by topic. For many open source or resource-constrained models, context windows limit the size of input for such tasks which can be constraining (Chung et al., 2024). Although context windows can be increased (Chen et al., 2023) and newer models are consistently increasing context lengths, this alone does not fix all problems as GLMs do not fully utilise long context windows (Liu et al., 2023) (Petroni et al., 2020).

1.2. Related Work

Previous methods of topic segmentation use either lexical or semantic similarity, and a variety of different machine learning approaches. However, there has been no research into the usage of LLMs for topic segmentation. LLMs have been shown to be effective at a variety of NLP tasks due to their vast general knowledge of language, a skill which is also required for topic segmentation.

2. Method

2.1. Datasets

There are 4 types of datasets used in the experiments in this work: a small human-annotated dataset, a scrape of English wikipedia, a 'concatenated' wikipedia scrape and a synthetic GPT3.5 generated dataset.

2.1.1. HUMAN-ANNOTATED DATASET

Lacking the resources to create a large gold standard dataset annotated by humans, this work uses a very small manually segmented dataset of 10 documents. These documents are a mix of news articles, wikipedia articles, and miscellaneous documents such as podcast transcripts and scientific reports. This was intended to represent varying difficulties of segmentation, and provide examples which could be manually inspected to interpret segmenter behaviour.

2.1.2. WIKIPEDIA DATASET

A plain text English wikipedia scrape [X] was used articles where headings were delimited by special characters. The

articles were then automatically segmented based on headings, and filtered to remove articles with very few segments, too short segments, or too muych punctuation such as tables and figures. This created idealised segments. After this process, there remained approximately 1000 segmented articles.

2.1.3. CONCATENATED WIKIPEDIA DATASET

We concatenated randomly sampled segments from the previous Wikipedia dataset in order to form new (incoherent) articles, with segments drawn from completely different domains. This should, intuitively, easier to segment.

2.1.4. SYNTHETIC DATASET

The final dataset used was generated synthetically by GPT-3.5. The source data was a mix of CTC sentinel data [X], UN-Peacekeeping corpus [X] and an internal dataset at Adarga. These documents were segmented by querying OpenAI's API with the model 'gpt-3.5-turbo-16k', using the overlapping prompt schema defined in section 2.3. While this dataset does not serve as a point of comparison between GLMs and other methods, as the ground truth is defined by a GLM, it was used to fine-tune an open-source LLM, as described in [X].

2.2. Evaluation

The methodology in evaluating the quality of a model's segmentation in this work follows Chris Fournier's *Evaluating Text Segmentation using Boundary Edit Distance* (Fournier, 2013). This work uses edited versions of the Boundary Edit Distance proposed in this work, along with the associated information recall metrics 'boundary precision' and 'boundary recall'.

The boundary edit distance algorithm pairs matches of segment boundaries between a reference segmentation and a hypothesised segmentation, regardless of distance between matches. Exact matches score 1 and boundaries without matches score 0, whilst matches within some distance n score partial points based on a weighting function. For this work, the function always decreases score linearly as the boundary gets farther away from the true boundary. Boundary Edit Distance (B) is the mean score for all these matches, while Boundary Precision/Recall (BP/BR) measure the proportion of the hypothesis/reference boundaries which are matched, respectively.

For further justification of the usage of the boundary edit distance algorithm rather than relying on more traditional metrics such as WD and Pk (Pevzner & Hearst, 2002), see *Evaluating Text Segmentation using Boundary Edit Distance* (Fournier, 2013), and our own investigations in segmenter evaluation metrics.

2.3. LLM-Based Text Segmentation

How can we get LLMs to output segment boundaries?

We might first consider passing in the input text as a prompt and asking the LLM to copy out the text, adding markers indicating where it has placed boundaries as is suggested by (Xing, 2024). However, not only is this wasteful of tokens, especially if the input text is long, but crucially, the GLM may fail to copy the input accurately, may change the formatting, and we found through qualitative experimentation that boundary placement was not any better than our final method. These problems are addressed by (Xing, 2024) through repeated prompting until the input and output sequence lengths match, but this still does not guarantee integrity of the data. In our use case, guaranteeing that the input data would remain the same was of the utmost importance, so we used a different strategy.

2.3.1. PROMPTING METHOD

We first annotate the text with indices between each sentence. For example 'Hello World. [1] The sky is blue. [2] The sun is is yellow. ' We then ask the GLM to return a list of indices corresponding to boundaries. In the previous example, the ideal response might be '1'. In practice, the texts and list of segment boundaries are much longer.

We add a system prompt which describes the segmentation task, desired output format and primes the model to be a talented linguist. We also add a variety of short examples in line with the few shot prompting technique (Brown et al., 2020), which increased performance. This did not necessarily reflect the fact that the GLM learnt how to segment better. Instead, through manual testing, we suspect that it learnt the ideal segment length and amount of information that should be contained within a segment, which was implicitly contained in the few shot prompt (and also the testing datasets, therefore increasing performance). This suggests that different implicit definitions of a segment could be imposed by a few shot prompt to a GLM, dependent upon use case.

An example prompt is contained in [XXX Appendix A].

2.3.2. OVERLAPPING PROMPTS

This prompting method works so long as the input text is within the context window of the LLM. At the time of experimentation, and with the use of gpt-3.5-turbo, we had a limit of 16k tokens. Many of our input texts exceeded this limit. Therefore, we needed to split up these long documents into smaller chunks that can be processed by the GLM. However, this cannot be done by simply splitting every at the nearest sentence before every 16k new tokens for two reasons. Firstly, we do not know whether this sentence boundary should serve as a segment boundary, and secondly, the GLM loses valuable context which helps to

choose where to place boundaries at the extremes of 16k tokens.

Therefore, we instead send 16k prompts that overlap. The overlap is calculated as twice maximum segment length. In our experiments, we set a maximum segment length of 750 tokens, thus there is an overlap of 1500 tokens between prompts. We must then decide which boundaries to accept in this overlapping region. Given two generations which were prompted by 1500 overlapping tokens, we choose to accept the segment boundaries contained within the first 750 tokens of the overlapping section from the first generation, and the boundaries in the final 750 tokens from the second prompt. We did not experiment with involving the responses from both outputs, but found that there seemed to be no sign of degrading performance towards segment boundaries.

While this method is wasteful of up to 1500 tokens per prompt, this is a small enough fraction of the 16k context that we were satisfied with the solution. If maximum segment lengths were much longer, say 5k tokens, the context may need to be limited more severely.

2.3.3. SEGMENT VALIDATION

We also performed some validation on the segments returned by the GLM. This primarily involved verify that the returned segments are within a maximum and minimum segment length. Segments that are too short (for example, a model would often return just a heading), were concatenated with another segment, and segments that are too long were recursively segmented by the same model, but with another prompt that asks the model to generate only one boundary at a time, which uses a similar few-shot prompting strategy to above. This way, a segment that is too long will be split in 2 recursively until all segments are within the specified lengths.

3. Experiments

3.1. Models

3.1.1. BASELINES

We use two naive baselines as a point of reference. First, a segmenter which splits every n sentences. We decided to split every 5 sentences which we call the *Split5Segmenter*. We also define a *RandomF0.1Segmenter* which splits at 10% of boundaries, placed randomly, with each potential boundary equally likely.

3.1.2. BERT SEGMENTER

Our existing method generates a sequence of sentence similarities using sentence embeddings generated by (Reimers & Gurevych, 2019). Similarities are calculated as a weighted sum of the cosine similarity to the previous n sentences.

Ideal boundaries are then generated as troughs in the sequence of similarities, before further processing to ensure there are no segments that are too long or too short, either in sentence length or in token length. We name this the *BERTSegmenter*. Further details are omitted for proprietary reasons.

3.1.3. BERT-GRAPH SEGMENTER

(Costacurta, 2023) describes 'text tiling' (topic segmentation) using BERT-generated similarity scores followed by graph clustering to find the best segments. This follows a similar methodology as the previous Exact details of this method can be found in the linked article. This model is called *BERTGraphSegmenter* in our experiments. Code for this model was copied from the repository linked in the article.

3.1.4. GPT-3.5

OpenAI's `gpt-3.5-turbo-16k` was queried using the prompting and segment validation strategy defined in Section 2. We call this model *GPT3.5*. We used the largest model available to us, which is the 16k token context window.

3.1.5. FLAN-T5-FINETUNED

Took Flan-T5 large [XX] and fine-tuned it a combination of wiki, concatenated wiki and synthetic data.

3.2. Quantitative Results

Due to resource constraints, we could not test *GPT3.5* on the full wikipedia or full concatenated wikipedia datasets. Instead, we took the largest subset that fit within resource constraints. We evaluated all other models on the full datasets to verify that similar results are obtained. Further details on the experimental procedure are witheld for proprietary reasons.

We evaluate primarily using the previously discussed boundary similarity metric with $n = 5$. We also looked at precision and recall which helped us characterize the behavior of each model. Results can be found in Table 1.

We find that ...

Results were also taken with $n = 2$ and $n = 10$. Full results for both the smaller and larger datasets and all evalutation metrics can be found in the Appendix A.

3.3. Qualitative Results

Through manual testing with the human-annotated dataset, we found that *GPT3.5* generally found the boundaries which seemed most reasonable from a human perspective, espe-

	Human	Wiki	Wiki-concat
GPT3.5	Row 1 Data	Row 1 Data	Row 1 Data
Flan-T5	Row 2 Data	Row 2 Data	Row 2 Data
BERTGraph	Row 3 Data	Row 3 Data	Row 3 Data
BERT	Row 4 Data	Row 4 Data	Row 4 Data
RandomF0.1	Row 5 Data	Row 5 Data	Row 5 Data
Split5	Row 6 Data	Row 6 Data	Row 6 Data

Table 1. Boundary similarity with $n = 5$ for the different models on the human-annotated, Wikipedia, and concatenated Wikipedia datasets.

cially for simple documents. The

BERT segmenters would find reasonable segments, but after the manual gluing and splitting procedure, would often lead to off-by-1 errors.

However, for documents with far more complex documents such as a podcast transcript, or with messy data like tables or artefacts from pdf to text conversions, *GPT3.5* would sometimes return indices with a regular pattern. For example, the GLM might return '$[1, 15, 22, ..., 76, 79, 82, 85, 88, ..., 184, 187, ...]$'. Often, the pattern would continue far beyond the number of sentences in the input indicating that the model became stuck in a regular pattern. Perhaps better prompt engineering, a more rigorous data-processing procedure or the use of newer models would help, but our current approach was resource constrained and required the ability to pass noisy documents to the model. A more thorough investigation of the logits computed by the model is required to understand how and when this occurs, and how to mitigate it.

4. Conclusion

We empirically compare Generative Large Language Models with previous methods using cosine similarities. We investigate different evaluation metrics for the problem of topic segmentation and

References

Brown, T. B., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., Agarwal, S., Herbert-Voss, A., Krueger, G., Henighan, T., Child, R., Ramesh, A., Ziegler, D. M., Wu, J., Winter, C., Hesse, C., Chen, M., Sigler, E., Litwin, M., Gray, S., Chess, B., Clark, J., Berner, C., McCandlish, S., Radford, A., Sutskever, I., and Amodei, D. Language models are few-shot learners, 2020. URL https://arxiv.org/abs/2005.14165.

Chen, S., Wong, S., Chen, L., and Tian, Y. Extending context window of large language models via positional interpolation, 2023. URL https://arxiv.org/abs/2306.15595.

Chung, H. W., Hou, L., Longpre, S., Zoph, B., Tay, Y., Fedus, W., Li, Y., Wang, X., Dehghani, M., Brahma, S., Webson, A., Gu, S. S., Dai, Z., Suzgun, M., Chen, X., Chowdhery, A., Castro-Ros, A., Pellat, M., Robinson, K., Valter, D., Narang, S., Mishra, G., Yu, A., Zhao, V. Y., Huang, Y., Dai, A. M., Yu, H., Petrov, S., Chi, E. H., Dean, J., Devlin, J., Roberts, A., Zhou, D., Le, Q. V., and Wei, J. Scaling instruction-finetuned language models. *J. Mach. Learn. Res.*, 25:70:1–70:53, 2024. URL https://jmlr.org/papers/v25/23-0870.html.

Costacurta, M. Text tiling done right: Building solid foundations for your personal llm, 2023.

Fournier, C. Evaluating text segmentation using boundary edit distance. In Schuetze, H., Fung, P., and Poesio, M. (eds.), *Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 1702–1712, Sofia, Bulgaria, August 2013. Association for Computational Linguistics. URL https://aclanthology.org/P13-1167.

Hearst, M. A. Texttiling: segmenting text into multi-paragraph subtopic passages. *Comput. Linguist.*, 23(1): 33–64, mar 1997. ISSN 0891-2017.

Liu, N. F., Lin, K., Hewitt, J., Paranjape, A., Bevilacqua, M., Petroni, F., and Liang, P. Lost in the middle: How language models use long contexts, 2023.

Petroni, F., Lewis, P., Piktus, A., Rocktäschel, T., Wu, Y., Miller, A. H., and Riedel, S. How context affects language models' factual predictions, 2020. URL https://arxiv.org/abs/2005.04611.

Pevzner, L. and Hearst, M. A. A critique and improvement of an evaluation metric for text segmentation. *Computational Linguistics*, 28(1):19–36, 2002. doi: 10.1162/089120102317341756. URL https://aclanthology.org/J02-1002.

Reimers, N. and Gurevych, I. Sentence-bert: Sentence embeddings using siamese bert-networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 11 2019. URL http://arxiv.org/abs/1908.10084.

Xing, L. *Versatile neural approaches to more accurate and robust topic segmentation.* PhD thesis, University of British Columbia, 2024. URL https://open.library.ubc.ca/collections/ubctheses/24/items/1.0440128.

A. Appendix